



Investor Presentation Q4 FY23

February 27, 2023

Except for the historical information contained herein, certain matters in this presentation including, but not limited to, statements as to: our financial position; our markets, market opportunity and growth drivers; AI adoption being at an inflection point; our expectation of strong quarter over quarter and accelerating year over year growth in Data Center; the Generative AI opportunity being significant and driving strong growth in Data Center; gamers responding enthusiastically to the new RTX 4090, 4080, and 4070 Ti desktop GPUs; the Gaming channel inventory correction being largely behind us and our expectation that the Professional Visualization correction will end in the first half of the year; interest in NVIDIA Omniverse continuing to build; our financial outlook, our expected tax rates and our expected capital expenditures for the first quarter of fiscal 2024; the benefits, impact, performance and availability of our products and technologies, including NVIDIA Ada Lovelace architecture, GeForce RTX 40 Series GPUs and laptops, NVIDIA H100 GPUs, DLSS 3, GeForce RTX GPUs, GeForce NOW, NVIDIA Reflex, NVIDIA DRIVE, NVIDIA DRIVE Orin, DRIVE Hyperion, NVIDIA NeMo LLM Service and NVIDIA BioNeMo LLM Service, NVIDIA Omniverse Computing Systems (OVX), NVIDIA Omniverse Cloud services, NVIDIA IGX platform, NVIDIA's acceleration stacks and ecosystems, NVIDIA Omniverse, NVIDIA DGX A100, NVIDIA Omniverse Avatar Cloud Engine, NVIDIA AI Enterprise, NVIDIA DRIVE Sim, Bluefield-3 DPU, Grace CPU Superchip, and NVIDIA DRIVE Thor; the benefits and impact of our collaborations or partnerships with Hyundai Motor Group, BYD, Polestar, Foxconn, Mercedes-Benz, Deutsche Bank, and Microsoft; accelerated computing being needed to tackle the most impactful opportunities of our time; expected TOPS processing; AI as the greatest technology force of our time; data centers across industries becoming AI factories; ; NVIDIA's AI expertise and scale helping to revolutionize businesses; Omniverse being essential for the next wave of AI – robotics; NVIDIA's value to every stakeholder in the ecosystem; the cost and time-to-solution savings of application speed-ups; our remaining repurchase authorization and dividend program plan; upcoming launches of our Data Center products; our Automotive design win pipeline, ramp and production expectations; NVIDIA accelerated computing being broadly recognized as the way to advance computing as Moore's law ends; the next wave of AI being robotics; building and operating Metaverse applications being the next wave; and our plan for 100% of our global electricity usage for our offices and data centers to be renewable by 2025 are forward-looking statements.

These forward-looking statements and any other forward-looking statements that go beyond historical facts that are made in this presentation are subject to risks and uncertainties that may cause actual results to differ materially. Important factors that could cause actual results to differ materially include: global economic conditions; our reliance on third parties to manufacture, assemble, package and test our products; the impact of technological development and competition; development of new products and technologies or enhancements to our existing product and technologies; market acceptance of our products or our partners' products; design, manufacturing or software defects; changes in consumer preferences and demands; changes in industry standards and interfaces; unexpected loss of performance of our products or technologies when integrated into systems and other factors.

NVIDIA has based these forward-looking statements largely on its current expectations and projections about future events and trends that it believes may affect its financial condition, results of operations, business strategy, short-term and long-term business operations and objectives, and financial needs. These forward-looking statements are subject to a number of risks and uncertainties, and you should not rely upon the forward-looking statements as predictions of future events. The future events and trends discussed in this presentation may not occur and actual results could differ materially and adversely from those anticipated or implied in the forward-looking statements. Although NVIDIA believes that the expectations reflected in the forward-looking statements are reasonable, the company cannot guarantee that future results, levels of activity, performance, achievements or events and circumstances reflected in the forward-looking statements will occur. Except as required by law, NVIDIA disclaims any obligation to update these forward-looking statements to reflect future events or circumstances. For a complete discussion of factors that could materially affect our financial results and operations, please refer to the reports we file from time to time with the SEC, including our most recent Annual Report on Form 10-K, Quarterly Reports on Form 10-Q, and Current Reports on Form 8-K. Copies of reports we file with the SEC are posted on our website and are available from NVIDIA without charge.

NVIDIA uses certain non-GAAP measures in this presentation including non-GAAP gross profit, non-GAAP gross margin, non-GAAP operating expenses, non-GAAP operating income, non-GAAP operating margin, non-GAAP net income, non-GAAP diluted earnings per share, and free cash flow. NVIDIA believes the presentation of its non-GAAP financial measures enhances investors' overall understanding of the company's historical financial performance. The presentation of the company's non-GAAP financial measures is not meant to be considered in isolation or as a substitute for the company's financial results prepared in accordance with GAAP, and the company's non-GAAP measures may be different from non-GAAP measures used by other companies. Further information relevant to the interpretation of non-GAAP financial measures, and reconciliations of these non-GAAP financial measures to the most comparable GAAP measures, may be found in the slide titled "Reconciliation of Non-GAAP to GAAP Financial Measures".





Content

- Q4 FY23 Earnings Summary
- Key Announcements This Quarter
- NVIDIA Overview
- Financials
- Reconciliation of Non-GAAP to GAAP Financial Measures


Q4 FY23 Earnings Summary



Highlights

- **Resumed sequential revenue growth driven by new Gaming product cycle; Y/Y decline reflects channel inventory corrections in Gaming and Pro Viz**
 - Total revenue down 21% Y/Y to \$6.05B, inline with outlook of \$6.00B +/- 2%
 - Data Center up 11% Y/Y to \$3.62B
 - Gaming down 46% Y/Y to \$1.83B
- **Data Center declined Q/Q on lower sales in China; expect strong Q/Q growth in Q1 and accelerating Y/Y growth throughout the year**
 - Q/Q decline driven by China reflecting COVID and other domestic issues; Y/Y growth driven by U.S. cloud service providers
 - New flagship H100 data center GPU revenue was already much higher than that of the A100, which declined sequentially
 - AI adoption is at an inflection point; the Generative AI opportunity is significant and driving strong growth in Data Center
- **Gaming grew Q/Q on strong reception of GeForce RTX 40 Series GPUs; down Y/Y due to channel inventory correction**
 - End demand in the seasonally strong fourth quarter was solid in most regions; early signs of recovery in China
 - Gamers are responding enthusiastically to the new RTX 4090, 4080, and 4070 Ti desktop GPUs
 - The channel inventory correction is largely behind us


Q4 FY23 Financial Summary



	GAAP			Non-GAAP		
	Q4 FY23	Y/Y	Q/Q	Q4 FY23	Y/Y	Q/Q
Revenue	\$6,051	-21%	+2%	\$6,051	-21%	+2%
Gross Margin	63.3%	-2.1 pts	+9.7 pts	66.1%	-0.9 pts	+10.0 pts
Operating Income	\$1,257	-58%	+109%	\$2,224	-40%	+45%
Net Income	\$1,414	-53%	+108%	\$2,174	-35%	+49%
Diluted EPS	\$0.57	-52%	+111%	\$0.88	-33%	+52%
Cash Flow from Ops	\$2,249	-26%	+474%	\$2,249	-26%	+474%

All dollar figures are in millions other than EPS. Refer to Appendix for reconciliation of Non-GAAP measures.


Data Center



Highlights

- Hyperscale customer revenue posted strong Q/Q growth
- Revenue growth from cloud service provider (CSP) customers last year significantly outpaced that of Data Center as a whole, as more enterprise customers move to a cloud-first approach
 - On a trailing four-quarter basis, CSPs drove about 40% of Data Center revenue
- Q/Q revenue decline was driven by lower sales into China, reflecting COVID and other domestic issues


Gaming



Highlights

- Q/Q growth was driven by the strong reception of our 40-series GeForce RTX GPUs based on the Ada Lovelace architecture
- Y/Y decline reflects the impact of the channel inventory correction, which is largely behind us
- The flagship RTX 4090 has quickly shot up in popularity on Steam to claim the top spot for the Ada architecture
- Over 400 games and applications support NVIDIA's RTX technology for real-time ray tracing and AI-powered graphics
- GeForce NOW cloud gaming service now has more than 25 million members in over 100 countries


Professional Visualization



Highlights

- Q/Q growth was driven by desktop workstations, with strength in the automotive and manufacturing industry verticals
- Y/Y decline reflects the impact of the channel inventory correction, which we expect to end in the first half of fiscal 2024
- Interest in NVIDIA Omniverse continues to build, with almost 300,000 downloads so far, and 185 connectors to third-party design applications


Automotive



Highlights

- Q/Q growth was driven primarily by AI Automotive Solutions. New program ramps at both Electric Vehicle and traditional OEM customers helped drive growth
- The NVIDIA DRIVE operating system received safety certification from TÜV SÜD, one of the most experienced and rigorous assessment bodies in the automotive industry

Sources & Uses of Cash



Highlights

- Y/Y decreases reflect lower operating income and timing of supplier payments and inventory deliveries, partially offset by lower accounts receivable due to strong collections
- Q/Q increase reflects timing of supplier payments and lower accounts receivable due to improved shipment linearity
- Returned \$1.15 billion to shareholders in the form of share repurchases and cash dividends
- Invested \$513M in capex (includes principal payments on PP&E)
- Ended the quarter with \$13.3B in gross cash and \$11.0B in debt; \$2.3B in net cash

Cash Flow from Operations (\$M)

Gross cash is defined as cash/cash equivalents & marketable securities.

Debt is defined as principal value of debt.


Net cash is defined as gross cash less debt.

Q1 FY24 Outlook

Revenue	\$6.5 billion , plus or minus 2% Expect sequential growth to be driven by each of our four major market platforms, led by strong growth in Data Center and Gaming
Gross Margins	64.1% GAAP and 66.5% non-GAAP, plus or minus 50 basis points
Operating Expense	Approximately \$2.53 billion GAAP and \$1.78 billion non-GAAP We plan to keep opex relatively flat at Q1 levels over the coming quarters
Other Income & Expense	Net income of approximately \$50 million for GAAP and non-GAAP Excluding gains and losses on non-affiliated investments
Tax Rate	13.0% GAAP and non-GAAP, plus or minus 1%, excluding discrete items
Capital Expenditures	Approximately \$350 million to \$400 million We expect capex to be in the range of \$1.1 to \$1.3 billion for fiscal year 2024

Refer to Appendix for reconciliation of Non-GAAP measures.

Key Announcements This Quarter



NVIDIA Ada Lovelace Supercharges 170+ Laptop Designs for Gamers and Creators

- GeForce RTX 40 Series laptops are up to 3x more power efficient than the previous generation and bring NVIDIA DLSS 3 and fifth-generation Max-Q technologies to laptops for the first time
- For gamers, RTX 40 Series laptops deliver up to 4x the performance in AAA games, and for creators, they can accelerate rendering scenes in creative apps such as Blender from 2.5 hours to just 10 minutes
- RTX 40 Series delivers enthusiast-class performance in laptops as slim as 14 inches, a new milestone
- The RTX 4090 and 4080 GPUs enable a new class of flagship laptops that are the world's fastest while new RTX 4070, 4060 and 4050 laptops are faster than the previous-generation flagship model
- RTX 40 Series flagship laptops start at \$1,999. RTX 4050 laptops start at \$999



GeForce RTX 4070 Ti Brings NVIDIA Ada Lovelace Architecture to \$799

- The new RTX 4070 Ti GPU is faster than the GeForce RTX 3090 Ti at nearly half the power
- For 3D creators, the RTX 4070 Ti delivers up to 70% increase in performance compared to the GeForce RTX 3070 Ti for popular tools like Blender, Unity, Unreal Engine and Chaos V-Ray
- RTX 4070 Ti is now available from top add-in card providers ASUS, Colorful, Gainward, GALAX, GIGABYTE, INNO3D, KFA2, MSI, Palit, PNY and ZOTAC, as well as from gaming system integrators and builders worldwide



GeForce NOW Coming to Cars

- GeForce NOW cloud gaming service will be coming to cars
- Hyundai Motor Group, BYD, and Polestar are the first automakers working with NVIDIA to deliver GeForce NOW in their vehicles
- Expands NVIDIA's vehicle infotainment offerings, which include a suite of products and services that improve the cockpit experience
- The new GeForce NOW offering can enhance time spent charging or riding in vehicles, as it enables front-seat occupants to stream games while parked, and passengers to game in the back seat if screens are available



New RTX 4080-Class Cloud Gaming Tier

- Announced a new, high-performance Ultimate membership tier of GeForce NOW, delivering GeForce RTX 4080-class graphics horsepower — nearly 1.75x over the previous generation
- Ultimate members will be the first to experience true PC gaming — streaming at up to 240 frames per second from the cloud with full ray tracing and DLSS 3
- With the addition of NVIDIA Reflex, GeForce NOW achieves click-to-pixel latency below 40 milliseconds — another first in cloud gaming
- Now available in North America and Europe, with continued rollout over the months to follow
- The GeForce NOW Ultimate membership is priced at \$19.99 per month or \$99.99 for six months



Foxconn Partners With NVIDIA to Build Automated Electric Vehicles

- Announced a strategic partnership with Hon Hai Technology Group (Foxconn) to develop automated and autonomous vehicle platforms
- As part of the agreement, Foxconn will be a tier-one manufacturer, producing electronic control units (ECUs) based on NVIDIA DRIVE Orin for the global automotive market
- Additionally, Foxconn-manufactured electric vehicles (EVs) will feature DRIVE Orin ECUs and DRIVE Hyperion sensors for highly automated driving capabilities



Mercedes-Benz Assembles Next-Gen Factories With NVIDIA Omniverse

- Mercedes-Benz announced that it is using the NVIDIA Omniverse platform to design and plan manufacturing and assembly facilities
- Mercedes-Benz planners can access the digital twin of the factory, reviewing and optimizing the plant as needed
- Allows Mercedes-Benz to synchronize plant locations anywhere in the world by connecting Omniverse with its in-house MO360 Data Platform
- The latest release of NVIDIA Omniverse Enterprise adds support for 4K real-time path tracing, new connectors — now over 150
- NVIDIA Omniverse Enterprise is available by subscription from BOXX, Dell, Z by HP and Lenovo, and channel partners including Arrow, ASK, PNY and Leadtek



NVIDIA Omniverse Avatar Cloud Engine (ACE) Now in Early Access

- ACE is a suite of cloud-native AI microservices that make it easier to build and deploy avatars, intelligent virtual assistants and digital humans at scale
- These AI assistants can be designed for organizations across industries, enabling organizations to enhance existing workflows and unlock new business opportunities
- ACE is one of several generative AI applications that will help creators accelerate the development of 3D worlds and the metaverse



Deutsche Bank, NVIDIA to Accelerate Adoption of AI

- Deutsche Bank (DB) announced a partnership with NVIDIA to accelerate the use of AI and machine learning in the financial services sector
- DB's initiatives to speed efforts to serve customers worldwide, develop new data-driven products and services, increase efficiency and recruit tech talent
- Together, DB and NVIDIA have initially focused on three potential implementations with a multi-year ambition to expand this to over a hundred
- With NVIDIA AI Enterprise software, DB's AI developers, data scientists and IT professionals will be able to build and run AI workflows anywhere, including in its hosted on-premises data centers and on Google Cloud, the bank's public cloud provider



Microsoft and NVIDIA Announce Expansive New Gaming Deal

- Microsoft and NVIDIA agreed to a 10-year partnership to bring Xbox PC games to the NVIDIA GeForce NOW cloud gaming service
- The agreement will enable gamers to stream Xbox PC titles from GeForce NOW to PCs, macOS, Chromebooks, smartphones and other devices
- It will also enable Activision Blizzard PC titles, such as Call of Duty, to be streamed on GeForce NOW after Microsoft's acquisition of Activision closes
- GeForce NOW eclipsed 1,500 games and 25 million members in over 100 countries



Mercedes Benz Taking Entire Vehicle Product Lifecycle Digital With NVIDIA AI and Omniverse

- NVIDIA founder and CEO Jensen Huang joined Mercedes-Benz CEO Ola Källenius on stage at the automaker's strategy update, showcasing progress in their landmark partnership
- The automaker unveiled its new operating system, MB.OS, a purpose-built, chip-to-cloud architecture that will be standard across its entire vehicle portfolio — delivering exceptional software capabilities and ease of use
 - At the heart of this architecture is NVIDIA DRIVE Orin, which delivers high-performance, energy-efficient AI compute to support a comprehensive sensor suite and software to safely enable enhanced assisted driving and, ultimately, level 3 conditionally automated driving
- Mercedes-Benz is also accelerating the development of these systems with the high-fidelity NVIDIA DRIVE Sim platform, built on NVIDIA Omniverse



NVIDIA Overview



NVIDIA pioneered accelerated computing to help solve impactful challenges classical computers cannot. A quarter of a century in the making, NVIDIA accelerated computing is broadly recognized as the way to advance computing as Moore's law ends and AI lifts off.

NVIDIA's platform is installed in several hundred million computers, is available in every cloud and from every server maker, powers 361 of the TOP500 supercomputers, and boasts over 3.5 million developers.

Headquarters: Santa Clara, CA




What Is Accelerated Computing?

A full-stack approach: silicon, systems, software

Not just a superfast chip – accelerated computing is a full-stack combination of:

- Chip(s) with specialized processors
- Algorithms in acceleration libraries
- Domain experts to refactor applications

To speed-up compute-intensive parts of an application.



Amdahl's law:

The overall system speed-up (S) gained by optimizing a single part of a system by a factor (s) is limited by the proportion of execution time of that part (p).

$$S = \frac{1}{(1 - p) + \frac{p}{s}}$$

For example:

- If 90% of the runtime can be accelerated by 100x, the application is sped up 9x
- If 99% of the runtime can be accelerated by 100x, the application is sped up 50x
- If 80% of the runtime can be accelerated by 500x, or even 1000x, the application is sped up 5x

Why Accelerated Computing?

Advancing computing in the post-Moore's Law era

Accelerated computing is needed to tackle the most impactful opportunities of our time—like AI, climate simulation, drug discovery, ray tracing, and robotics.


NVIDIA is uniquely dedicated to accelerated computing—working top-to-bottom—refactoring applications and creating new algorithms, and bottom-to-top—inventing new specialized processors, like RT Core and Tensor Core.

“It’s the end of Moore’s Law as we know it.”

- John Hennessy Oct 23, 2018

“Moore’s Law is dead.”

- Jensen Huang, GTC 2013



AI Is the Greatest Technology Force of Our Time

Data centers across industries will become AI factories

AI has fundamentally changed what software can make and how you make software.

Companies are processing & refining their data, making AI software—becoming intelligence manufacturers. Their data centers are AI factories.

The first wave of AI is learned perception and inference, like recognizing images, understanding speech, recommending a video, or an item to buy.

The next wave of AI is robotics—AI planning actions. Digital robots, avatars, and physical robots will perceive, plan and act.

NVIDIA's acceleration stacks and ecosystems help bring AI to the world's largest industries.


NVIDIA's world-class AI expertise and scale can help revolutionize businesses.



Contact Center AI
500M Calls / Day




Meeting Transcription
3B Meeting Minutes / Day



Public Safety
>1B Smart City Cameras Deployed



Retail Asset Protection
\$94.5B Inventory Loss / Year



Medical Imaging
10M Diagnostic Scans / Day




Industrial Inspection
\$32M Vision Sensors Installed by 2025



Transportation
10T Miles / Year



Credit Card Fraud
1.28B Credit Transactions / Day



Product Recommendations
1B E-Commerce Visitors / Day

Source: Nilson Report, IHS Markit, Similar Web, NRF, WHO, ABI and NVIDIA internal analysis

Building and Operating Metaverse Applications Is the Next Wave

NVIDIA Omniverse—Runs on NVIDIA OVX servers | RTX workstations | Enterprise software | Cloud services

NVIDIA Omniverse is a software platform for building and operating metaverse applications.


Our initial focus is on industrial metaverses, such as digital twins used to emulate the behavior of products or factories in the physical world.

Omniverse uses a real-time, large-scale 3D database that connects to 3D worlds via the USD (Universal Scene Descriptor) framework.

Just as the internet connects websites over HTML, Omniverse connects 3D worlds over USD.

Omniverse is essential for the next wave of AI—robotics—where AI interacts with the physical world.

Applications built to run on Omniverse are like portals into the Omniverse virtual world.



Product Development


Connecting 3D Creators & AI Assistants in Virtual Worlds

Operate Digital Twin

Connecting Robots in a Virtual World

NVIDIA's Accelerated Computing Platform

Full-stack innovation across silicon, systems and software



With nearly three decades of a singular focus, NVIDIA is expert at accelerating software and scaling compute by a Million-X, going well beyond Moore's law.

Accelerated computing is a full-stack challenge, demanding deep understanding of the problem domain, optimizing across every layer of computing, and all three chips —GPU, CPU, and DPU.


Scaling across multi-GPUs and multi-nodes is a data center-scale challenge and requires treating the network and storage as part of the computing fabric.

Our platform extends from PCs to supercomputing centers, enterprise data centers, cloud and edge environments.

NVIDIA's Multi-Sided Platform and Flywheel

NVIDIA is valued by every stakeholder in the ecosystem:


- **For developers** – NVIDIA's One Architecture and large installed base give developer's software the best performance and greatest reach
- **For computer makers and CSPs** – NVIDIA's rich suite of Acceleration Platforms lets partners build one offering to address large markets including media & entertainment, healthcare, transportation, energy, financial services, manufacturing, retail, and more
- **For customers** – NVIDIA is offered by virtually every computing provider and accelerates the most impactful applications from cloud to edge
- **For NVIDIA** – Deep engagement with developers, computing providers, and customers in diverse industries enables unmatched expertise, scale, and speed of innovation across the entire accelerated computing stack – propelling the flywheel




Full-Stack & Data Center Scale Acceleration

Drive Significant Cost Savings and Workload Scaling

Classical Computing—92 CPU-only servers
\$3.3M (including switches, cables, racks)



Accelerated Computing—1 NVIDIA DGX A100
\$220,000 DGX and \$100,000 NVIDIA AI software



**10X lower cost
14X better energy-efficiency**

Cost comparison example based on latest available NVIDIA A100 GPU and Intel CPU inference results in the commercially available category of the MLPerf industry benchmark; includes related infrastructure costs such as networking.

New NVIDIA Software and Services

Enabling the World's Enterprises to Revolutionize Industries with AI

NVIDIA AI Enterprise

The operating engine of AI for end-to-end data-driven software development.

One engine license accelerates end-to-end modern AI and data science.

One engine license unlocks wealth of data processing, AI, and robotics frameworks and applications—e.g., RAPIDS, Spark, Merlin, Monai, Metropolis, cuOpt, Morpheus, Tokkio.

Per GPU On-Prem Subscription

Per GPU-HR Cloud Consumption

NVIDIA Omniverse

A platform for designing, building, and operating 3D and virtual world simulations.

Consists of a virtual world engine, USD connectors, and portals browsing the virtual world simulation.

Omniverse is an enterprise application that connects architects, designers, hardware and software engineers, marketers, to supply-chain and factory planners.

Per Connection On-Prem Subscription

Per GPU-HR Cloud Consumption

NVIDIA Nemo LLM

NVIDIA-hosted cloud service for training Large Language Models to perform specific tasks—e.g., summarize legal documents, write marketing copy, analyze market sentiment, chatbot to support customers, search documents, write and document code, paraphrase

Nemo can help thousands of companies, train language AI's to do hundreds of tasks, in 10's of languages.

Per GPU On-Prem Subscription

Per GPU-HR Cloud Consumption

NVIDIA BioNemo

NVIDIA-hosted cloud service for training and deploying large biomolecular models that understand the language of chemistry, proteins, RNA, and DNA.

BioNemo can help researchers, biotech, and pharma companies to process chemical and biological datasets to accelerate drug discovery.

Per GPU On-Prem Subscription

Per GPU-HR Cloud Consumption

Giant Market Opportunity



Gaming
Over 3B gamers and creators, a quarter of them spending over \$100/year for GPUs in desktops, laptops, cloud or consoles


NVIDIA AI Enterprise Software
50M enterprise server installed base; per-server, per-year subscription price

Omniverse Enterprise Software
Over 45M designers and creators; 10s of millions of digital twins—per-user/digital twin, per-year subscription price



Chips and Systems
~20M servers/year—GPUs, CPUs, DPUs, NICs, switches

Automotive
100M vehicles/year hardware opportunity; 100s of millions of AV vehicles installed base software opportunity

\$1 Trillion Opportunity



Driving Strong & Profitable Growth



NVIDIA Gross Margins Reflect Value of Acceleration


Accelerated computing requires full-stack and data center-scale innovation across silicon, systems, algorithms and applications.

Significant expertise and effort are required, but application speed-ups can be incredible, resulting in dramatic cost and time-to-solution savings.

For example, 10 NVIDIA HGX nodes with 80 NVIDIA A100 GPUs that cost \$4M can replace 920 nodes of CPU servers that cost over \$50M for AI inference.

NVIDIA chips carry the value of the full-stack, not just the chip.

Cost comparison example based on latest available NVIDIA A100 GPU and Intel CPU inference results in the commercially available category of the MLPerf industry benchmark; includes related infrastructure costs such as networking.



FY23 financial metrics reflect a \$2.2B charge for inventory and related reserves primarily related to Data Center and Gaming. Fiscal year ends in January. Refer to Appendix for reconciliation of Non-GAAP measures. Gross margins are rounded to the nearest percent.

Strong Cash Flow Generation

Free Cash Flow (Non-GAAP)



Capital Allocation

Share Repurchase

Resumed Buybacks in Q1 FY 2023
\$10.0B repurchased in FY23; \$7.2B Remaining Authorization Through Dec 2023 as of Jan 29, 2023

Dividend

\$398M in FY 2023
Plan to Maintain¹


Strategic Investments

Growing Our Talent
Platform Reach & Ecosystem

Fiscal year ends in January. Refer to Appendix for reconciliation of Non-GAAP measures.

¹ Subject to continuing determination by our Board of Directors.

Our Market Platforms at a Glance



Data Center

56% of FY23 revenue

FY23 Revenue \$15.0B
5-yr CAGR 51%

DGX/HGX/EGX/IGX systems
GPU | CPU | DPU | Networking
NVIDIA AI software

Gaming

33% of FY23 revenue

FY23 Revenue \$9.1B
5-yr CAGR 10%

GeForce GPUs for PC gaming
GeForce NOW cloud gaming

Professional Visualization

6% of FY23 revenue

FY23 Revenue \$1.5B
5-yr CAGR 11%

Quadro/NVIDIA RTX GPUs
for workstations
Omniverse software

Automotive



3% of FY23 revenue

FY23 Revenue \$0.9B
5-yr CAGR 10%

DRIVE Hyperion sensor architecture
with AGX compute
DRIVE AV & IX full stack software
for ADAS, AV & AI cockpit


Data Center

The leading computing platform for AI, HPC & graphics




Data Center


Strong growth fueled by AI performance leadership and huge developer ecosystem



Accelerating Adoption with
Every Architecture




#1 in AI Training &
Inference Performance



Modern AI is a Data Center Scale Computing Workload


Data centers are becoming AI factories: data as input, intelligence as output



Wave of New Data Center Products

Ramping new architectures for GPU, CPU and DPU


H100 GPU



World's Most Advanced Chip
80B Transistors
Transformer Engine – 6X Perf
Confidential Computing
4th Gen NVLink—7X PCIe Gen5

2H FY23

Bluefield-3 DPU



First 400 Gb/s DPU
Line-rate processing of software-defined networking, storage, and cybersecurity
VMware vSphere 8 integration
Zero-trust security
~600 infrastructure software partners

1H FY23

Grace CPU Superchip




High Performance CPU for HPC and AI
144 Cores | 740 SPECrate[®]2017_int_base est.
1TB/s Memory Bandwidth
2X Perf/Watt Over Traditional Servers
Runs NVIDIA Computing Stacks

1H FY24

Gaming

GeForce—the world's largest gaming platform



Leader in PC Gaming




Strong #1 market position with over 80% share
15 of the Top 15 most popular GPUs on Steam
Leading performance & innovation
200M+ gamers on GeForce

Growth Drivers

Rising adoption of NVIDIA RTX
Expanding universe of gamers & creators
Gaming laptops & game consoles
GeForce NOW Cloud gaming

Strong Gaming Fundamentals

New generation, more gamers




New generation, more gamers

Expanding universe of gamers and creators

Robust NVIDIA ecosystem

Professional Visualization

Workstation graphics



Leader in Workstation Graphics


90%+ market share in graphics for workstations
45M Designers and Creators
Strong software ecosystem with over 100 supported applications

Growth Drivers

Ray Tracing and AI revolutionizing design
Expanding universe of designers and creators
Collaborative 3D design / Omniverse
Hybrid work environments

Automotive

Autonomous Vehicles (AV) & AI Cockpit




Leader in Autonomous Driving
Historical revenue driven largely by infotainment
Future growth primarily fueled by NVIDIA DRIVE, our AV and
AI cockpit platform with full software stack
Over \$11B design win pipeline through FY28 based on DRIVE
Orin, which started ramp in FY23
Next-generation DRIVE Thor to ramp in FY26

Growth Drivers
Adoption of centralized car computing and
software-defined vehicle architectures
AV software and services:
Mercedes Benz FY25 SOP*
Jaguar Land Rover FY26 SOP


*SOP = Start of Production

Financials




Annual Cash & Cash Flow Metrics


Operating Income (Non-GAAP) — \$M




Operating Cash Flow — \$M



Free Cash Flow (Non-GAAP) — \$M



Cash Balance — \$M



*Cash balance is defined as cash and cash equivalents plus marketable securities
Refer to Appendix for reconciliation of non-GAAP measures*

Corporate Responsibility


Environmentally Conscious



NVIDIA's two HQ campuses have received **LEED Gold status**



23 of Top 30 Supercomputers on the Nov 2022 Green500 powered by NVIDIA including the #1 system, Henri



We Plan For **100% of Our Global Electricity Usage For Our Offices and Data Centers** to Be Renewable by 2025

A Place For People To Do Their Life's Work



"100 Best Companies to Work For"
FORTUNE

"America's Most Just Companies"
FORBES

"Most Responsible Companies"
NEWSWEEK

"Best Places to Work for LGBT Equality"
HUMAN RIGHTS CAMPAIGN

Management

Time Magazine's 100 Most Influential Companies

CEO Magazine's 10 Best CEOs

Fortune's World's Most Admired Companies

Wall Street Journal's Management Top 250 All-Stars

Corporate Governance

38% Of Board is Gender, Racially, or Ethnically Diverse

92% of Directors are Independent

Reconciliation of Non-GAAP to GAAP Financial Measures



Reconciliation of Non-GAAP to GAAP Financial Measures

	Non-GAAP	Acquisition Termination Cost	Acquisition-Related and Other Costs (A)	Stock-Based Compensation (B)	IP-Related Cost	Other (C)	Tax Impact of Adjustments	GAAP
Q4 FY23								
Gross margin (\$ in million)	\$3,999	—	(120)	(30)	(16)	—	—	\$3,833
	66.1%	—	(2.0)	(0.5)	(0.3)	—	—	63.3%
Operating income (\$ in million)	\$2,224	—	(174)	(739)	(16)	(38)	—	\$1,257
Net income (\$ in million)	\$2,174	—	(174)	(739)	(16)	(49)	218	\$1,414
Shares used in diluted per share calculation (millions)	2,477	—	—	—	—	—	—	2,477
Diluted EPS	\$0.88	—	—	—	—	—	—	\$0.57

A. Consists of amortization of intangible assets, transaction costs, and certain compensation charges.

B. Stock-based compensation charge was allocated to cost of goods sold, research and development expense, and sales, general and administrative expense.

C. Other comprises of restructuring costs and other and net losses from non-affiliated investments



Reconciliation of Non-GAAP to GAAP Financial Measures (contd)

Gross Margin	Non-GAAP	Acquisition-Related and Other Costs (A)	Stock-Based Compensation (B)	GAAP
Q4 FY2022	67.0%	(1.1)	(0.5)	65.4%
Q1 FY2023	67.1%	(1.1)	(0.5)	65.5%
Q2 FY2023	45.9%	(1.8)	(0.6)	43.5%
Q3 FY2023	56.1%	(2.0)	(0.5)	53.6%

A. Consists of amortization of intangible assets

B. Stock-based compensation charge was allocated to cost of goods sold

Reconciliation of Non-GAAP to GAAP Financial Measures (contd.)

Gross Margin (\$ in Millions & Margin Percentage)	Non-GAAP	Acquisition-Related and Other Costs (A)	Stock-Based Compensation (B)	IP-Related Costs	GAAP
FY 2019	\$7,233	—	(27)	(35)	\$7,171
	61.7%	—	(0.2)	(0.3)	61.2%
FY 2020	\$6,821	—	(39)	(14)	\$6,768
	62.5%	—	(0.4)	(0.1)	62.0%
FY 2021	\$10,947	(425)	(88)	(38)	\$10,396
	65.6%	(2.6)	(0.5)	(0.2)	62.3%
FY 2022	\$17,969	(344)	(141)	(9)	\$17,475
	66.8%	(1.4)	(0.5)	—	64.9%
FY 2023	\$15,965	(455)	(138)	(16)	\$15,356
	59.2%	(1.7)	(0.5)	(0.1)	56.9%

A. Consists of amortization of intangible assets and inventory step-up

B. Stock-based compensation charge was allocated to cost of goods sold



Reconciliation of Non-GAAP to GAAP Financial Measures (contd.)

Operating Margin (\$ in Millions & Margin Percentage)	Non-GAAP	Acquisition Termination Cost	Acquisition- Related and Other Costs (A)	Stock-Based Compensation (B)	IP-Related Cost	Other (C)	GAAP
FY 2019	\$4,407	—	(2)	(557)	(35)	(9)	\$3,804
	37.6%	—	—	(4.7)	(0.3)	(0.1)	32.5%
FY 2020	\$3,735	—	(31)	(844)	(14)	—	\$2,846
	34.2%	—	(0.3)	(7.7)	(0.1)	—	26.1%
FY 2021	\$6,803	—	(836)	(1,397)	(38)	—	\$4,532
	40.8%	—	(5.0)	(8.4)	(0.2)	—	27.2%
FY 2022	\$12,690	—	(636)	(2,004)	(9)	—	\$10,041
	47.2%	—	(2.5)	(7.4)	—	—	37.3%
FY 2023	\$9,040	(1,353)	(674)	(2,710)	(16)	(63)	\$4,224
	33.5%	(5.0)	(2.5)	(10.0)	(0.1)	(0.2)	15.7%

A. Consists of amortization of acquisition-related intangible assets, inventory step-up, transaction costs, compensation charges, and other costs

B. Stock-based compensation charge was allocated to cost of goods sold, research and development expense, and sales, general and administrative expense

C. Comprises of legal settlement costs, contributions, and restructuring and other charges



Reconciliation of Non-GAAP to GAAP Financial Measures (contd.)

(\$ in Millions)	Free Cash Flow	Purchases Related to Property and Equipment and Intangible Assets	Principal Payments on Property and Equipment and Intangible Assets	Net Cash Provided by Operating Activities
FY 2019	\$3,143	600	—	\$3,743
FY 2020	\$4,272	489	—	\$4,761
FY 2021	\$4,677	1,128	17	\$5,822
FY 2022	\$8,049	976	83	\$9,108
FY 2023	\$3,750	1,833	58	\$5,641

Reconciliation of Non-GAAP to GAAP Financial Measures

(\$ in Millions)	Q1 FY24 Outlook
Non-GAAP gross margin	66.5%
Impact of stock-based compensation expense, acquisition-related costs, and other costs	(2.4%)
GAAP gross margin	64.1%
Non-GAAP operating expenses	\$1,775
Impact of stock-based compensation expense and acquisition-related costs	750
GAAP operating expenses	\$2,525

